

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report, accuracy_score, precision_score, recall_score, f1_score

In [2]: data = pd.read_csv('../input/Social_Network_Ads2.csv')
data.head(5)
```

```
-----
FileNotFoundError                                Traceback (most recent call last)
/tmp/ipykernel_6456/1954719527.py in <module>
----> 1 data = pd.read_csv('../input/Social_Network_Ads2.csv')
      2 data.head(5)

~/local/lib/python3.10/site-packages/pandas/util/_decorators.py in wrapper(*args, **kwargs)
    209     else:
    210         kwargs[new_arg_name] = new_arg_value
--> 211     return func(*args, **kwargs)
    212
    213     return cast(F, wrapper)

~/local/lib/python3.10/site-packages/pandas/util/_decorators.py in wrapper(*args, **kwargs)
    329         stacklevel=find_stack_level(),
    330     )
--> 331     return func(*args, **kwargs)
    332
    333     # error: "Callable[[VarArg(Any), KwArg(Any)], Any]" has no
    334
~/local/lib/python3.10/site-packages/pandas/io/parsers/readers.py in read_csv(filepath_or_buffer, sep, delimiter, header, names, index_col, usecols, squeeze, prefix, mangle_dupe_cols, dtype, engine, converters, true_values, false_values, skipinitialspace, skiprows, skipfooter, nrows, na_values, keep_default_na, na_filter, verbose, skip_blank_lines, parse_dates, infer_datetime_format, keep_date_col, date_parser, dayfirst, cache_dates, iterator, chunksize, compression, thousands, decimal, lineterminator, quotechar, quoting, doublequote, escapechar, comment, encoding, encoding_errors, warn_bad_lines, warn_bad_lines, delim_whitespace, low_memory, memory_map, float_precision, storage_options)
    948     kws.update(kws_defaults)
    949
--> 950     return _read(filepath_or_buffer, kwds)
    951
    952

~/local/lib/python3.10/site-packages/pandas/io/parsers/readers.py in _read(filepath_or_buffer, kwds)
    603
    604     # Create the parser.
--> 605     parser = TextFileReader(filepath_or_buffer, **kwds)
    606
    607     if chunksize or iterator:
    608
~/local/lib/python3.10/site-packages/pandas/io/parsers/readers.py in __init__(self, f, engine, **kwds)
   1440
   1441     self.handles: IOHandles | None = None
--> 1442     self._engine = self._make_engine(f, self.engine)
   1443
   1444     def close(self) -> None:

~/local/lib/python3.10/site-packages/pandas/io/parsers/readers.py in _make_engine(self, f, engine)
   1733         if "b" not in mode:
   1734             mode += "b"
--> 1735         self.handles = get_handle(
   1736             f,
   1737             mode,

~/local/lib/python3.10/site-packages/pandas/io/common.py in get_handle(path_or_buf, mode, encoding, compression, memory_map, is_text, errors, storage_options)
    854     if ioargs.encoding and "b" not in ioargs.mode:
    855         # Encoding
--> 856         handle = open(
    857             handle,
    858             ioargs.mode,

FileNotFoundError: [Errno 2] No such file or directory: '../input/Social_Network_Ads2.csv'
```

```
In [3]: data = pd.read_csv('Social_Network_Ads2.csv')
data.head(5)
```

```
Out[3]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	15000	0
1	15610944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [4]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
 #   Column          Non-Null Count  Dtype
---  --
 0   User ID         400 non-null    int64
 1   Gender          400 non-null    object
 2   Age             400 non-null    int64
 3   EstimatedSalary 400 non-null    int64
 4   Purchased       400 non-null    int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
In [5]: data.describe()
```

```
Out[5]:
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.478954
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
In [6]: X = dataset.iloc[:, [2, 3]].values
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_6456/3223724410.py in <module>
----> 1 X = dataset.iloc[:, [2, 3]].values

NameError: name 'dataset' is not defined
```

```
In [7]: X = data.iloc[:, [2, 3]].values
```

```
In [8]: y = data.iloc[:, 4].values
```

```
print(X[:3, :])
print("-"*15)
print(y[:3])

[[ 19 15000]
 [ 35 20000]
 [ 26 43000]]
-----
[0 0 0]
```

```
In [9]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)
```

```
In [10]: print(X_train[:3])
print("-"*15)
print(y_train[:3])
print("-"*15)
print(X_test[:3])
print("-"*15)
print(y_test[:3])
```

```
[[ 44 39000]
 [ 32 120000]
 [ 38 50000]]
-----
[0 1 0]

[[ 30 87000]
 [ 38 50000]
 [ 35 76000]]
-----
[0 0 0]
```

```
In [11]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(300, 2)
(100, 2)
(300,)
(100,)
```

```
In [12]: from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)
```

```
In [13]: print(X_train[:3])
print("-"*15)
print(X_test[:3])
```

```
[[ 0.28194044 -0.80670699]
 [-0.60873761  1.46173768]
 [-0.01254409 -0.5677824 ]]
-----
[[ -0.80480212  0.50496393]
 [-0.01254409 -0.5677824 ]
 [-0.01254409 -0.5677824 ]]
[-0.30964085  0.1570462 ]]
```

```
In [14]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0, solver='lbfgs')
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)


print(X_test[:10])
print("-"*15)
print(y_pred[:10])
```

```
[[ -0.80480212  0.50496393]
 [-0.01254409 -0.5677824 ]
 [-0.30964085  0.1570462 ]
 [-0.80480212  0.27391877]
 [-0.30964085 -0.5677824 ]
 [-1.10309888 -1.43757073]
 [-0.70576986 -1.58254245]
 [-0.21060859  2.15757314]
 [-1.95318916 -0.04596581]
 [ 0.8787462  -0.77073441]]
-----
[0 0 0 0 0 0 0 1 0 1]
```

```
In [15]: from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred)
print("Confusion matrix : \n",cm)

Confusion matrix :
[[65  3]
 [ 0 24]]
```

```
In [16]: import seaborn as sns
import matplotlib.pyplot as plt
sns.heatmap(cm,annot=True)
plt.show()
```



```
In [17]: from sklearn.metrics import accuracy_score
print("Accuracy is :", accuracy_score(y_test,y_pred)*100,'%')
```

```
Accuracy is : 98.0 %

File ~/tmp/ipykernel_6456/3263289684.py, line 4
      Accuracy is : 98.0 %
      ^
SyntaxError: invalid syntax
```

```
In [18]: from sklearn.metrics import accuracy_score
print("Accuracy is :", accuracy_score(y_test,y_pred)*100,'%')
```

```
Accuracy is : 89.0 %
```

```
In [19]: from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
In [20]: precision=precision_score(y_test,y_pred)
print('Precision: %f' % precision)
```

```
Precision: 0.888889
```

```
In [21]: recall=recall_score(y_test,y_pred)
print('Recall: %f' % recall)
```

```
Recall: 0.750000
```

```
In [22]: f1=f1_score(y_test,y_pred)
print('F1 score: %f' % f1)
```

```
F1 score: 0.813559
```


```
In [23]: print(y_pred[:20])
print(y_test[:20])
```

```
[0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 1 0]
[0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0]
```

```
In [24]: from matplotlib.colors import ListedColormap
X_set, y_set = X_train, y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
             alpha = 0.6, cmap = ListedColormap(('red', 'green'))))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Logistic Regression (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
```

"c" argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use the "color" keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.


"c" argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use the "color" keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.



```
In [25]: # Visualizing the Test set results
from matplotlib.colors import ListedColormap
X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
             alpha = 0.6, cmap = ListedColormap(('red', 'green'))))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Logistic Regression (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

"c" argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use the "color" keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.

"c" argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use the "color" keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.



In [ ]: